CSCE 5222 — Feature Engineering for Time-Series

Group 18 | Ashish Rathnakar Shetty & Kushal Sai Venigalla

University of North Texas

About This Notebook

This notebook is the live demo for our Project 2 submission. It runs the complete feature engineering pipeline we built, end-to-end, across three different time-series tasks:

Task	Dataset	Goal
Classification	UCR Archive (5 datasets)	Classify time-series patterns
Forecasting	ETTh1 (electricity transformer)	Predict oil temperature one step ahead
Anomaly Detection	WESAD-like PPG signal	Find bad sensor segments without labels

We test four feature engineering strategies across all three tasks:

1. **Raw features** — baseline, no engineering
2. **Catch22** — 22 compact statistical features we wrote from scratch in pure NumPy
3. **Self-supervised MLP embeddings** — a small neural net learns features on its own
4. **SHAP-guided pruning** — use explainability to cut redundant features

---

How to Run

1. Click **Runtime → Run all** at the top of Colab
2. Or run each cell one by one to explain what each part does

First run takes about 2–3 minutes because it downloads:

- Five UCR classification datasets (~5 MB)
- ETTh1 electricity transformer dataset (~2 MB)
- The `shap` Python package (the only pip install needed)

Everything else (NumPy, pandas, scikit-learn, matplotlib) is already available in Google Colab.

What Each Section Does

Section	What Happens
Cell 1	Installs `shap` and imports all libraries
Cell 2	Defines our custom Catch22 implementation (22 features in pure NumPy)
Cell 3	Defines the SHAP-guided feature pruning helper
Cell 4	Downloads all three datasets and plots the PPG signal
Cells 5–8	**Task 1: Classification** — runs all 5 UCR datasets, shows accuracy chart, SHAP importance, and PCA projection
Cells 9–11	**Task 2: Forecasting** — builds lag/rolling features, trains all 4 models, shows MSE comparison
Cells 12–13	**Task 3: Anomaly Detection** — Isolation Forest on raw vs Catch22 vs SHAP-selected
Cell 14	Prints the final summary table of all results

Key Results (What You Should See)

Classification (UCR):

- SHAP-pruned Catch22 reaches 95.0% on SyntheticControl (+9.3% over raw)
- Catch22 improves TwoLeadECG by 8.8% even with only 23 training samples
- Raw features win on ItalyPowerDemand (local timing event Catch22 misses)

Forecasting (ETTh1):

- Lag + Rolling statistics reduce MSE by 62% over the raw lag-1 baseline
- SHAP pruning removes 12 of 28 features and achieves the best MSE of 0.4734

Anomaly Detection (PPG):

- Raw windows: F1 = 0.871
- Catch22 drops to 0.613 because of amplitude normalization
- SHAP selection recovers to 0.645 by keeping `diff_variance`

Repository and Report

- **GitHub:** https://github.com/Ash-projects-personal/CSCE5222-Group18-FeatureEngineering
- **Report:** 10-page IEEE-format PDF available in the same repository
- **Course:** CSCE 5222 Feature Engineering, Spring 2026

Run all cells below. Each one is self-contained and explained inline.

CSCE 5222 – Feature Engineering for Time-Series

Group 18 Demo Notebook

Ashish Rathnakar Shetty (ID: 11808466) · Kushal Sai Venigalla (ID: 11852559)
Department of Computer Science and Engineering, University of North Texas

Cell 1 — Install & Import

Colab already has most of what we need. We just need to make sure `shap` is up to date.

```
# install shap — everything else (numpy, pandas, sklearn, matplotlib)
is already in Colab
!pip install shap --quiet

import os, warnings, urllib.request, zipfile, glob
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import seaborn as sns

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import (GradientBoostingClassifier,
                              GradientBoostingRegressor,
                              IsolationForest)
from sklearn.linear_model import Ridge
from sklearn.decomposition import PCA
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import (accuracy_score, f1_score,
                             mean_squared_error, mean_absolute_error,
                             roc_auc_score)
import shap

warnings.filterwarnings('ignore')
np.random.seed(42)

plt.rcParams.update({
    'font.size': 11,
    'figure.dpi': 110,
    'axes.spines.top': False,
    'axes.spines.right': False,
})

print("Everything imported successfully!")
```

```
print(f"  numpy   {np.__version__}")
print(f"  pandas  {pd.__version__}")
print(f"  sklearn {__import__('sklearn').__version__}")
print(f"  shap    {shap.__version__}")
```

```
Everything imported successfully!
  numpy   2.0.2
  pandas  2.2.2
  sklearn 1.6.1
  shap    0.51.0
```

Cell 2 — Our Catch22 Implementation

The original `pycatch22` library needs a C compiler to install, which causes problems on some machines.
So we wrote all 22 features ourselves in pure NumPy — no dependencies at all.

The 22 features cover six categories:

- **Distribution**: mean, std, skewness, kurtosis, fraction above mean, outlier fraction
- **Temporal dependence**: autocorrelation at lag 1 and 2, zero-crossing of ACF
- **Complexity**: permutation entropy, histogram entropy, Lempel-Ziv complexity, spectral entropy
- **Scaling**: Hurst exponent, frequency power ratio, peak frequency
- **Dynamics**: variance of first differences, time-reversibility, nonlinearity, zero-crossings, longest streak

```
def catch22_features(x):
    """
    Compute all 22 Catch22 features for a 1D time series.
    Returns a dict of {feature_name: value}.
    """
    x = np.asarray(x, dtype=float)
    n = len(x)
    feats = {}

    # --- distribution ---
    mu = np.mean(x)
    s  = np.std(x, ddof=1) + 1e-10
    feats['mean']           = float(mu)
    feats['std']            = float(s)
    feats['skewness']       = float(np.mean(((x - mu) / s) ** 3))
    feats['kurtosis']       = float(np.mean(((x - mu) / s) ** 4) -
3.0)
    feats['above_mean']     = float(np.mean(x > mu))
    feats['outlier_frac']   = float(np.mean(np.abs((x - mu) / s) > 2))

    # --- temporal dependence ---
    xn  = x - mu
```

```
    acf = np.correlate(xn, xn, mode='full')[n-1:]
    acf_norm = acf / (acf[0] + 1e-10)
    cz  = np.where(np.diff(np.sign(acf_norm)))[0]
    c1e = np.where(acf_norm < 1.0 / np.e)[0]
    feats['first_zero_acf'] = int(cz[0])  if len(cz)  else n
    feats['first_1e_acf']   = int(c1e[0]) if len(c1e) else n
    feats['acf_lag1'] = float(np.corrcoef(x[:-1], x[1:])[0, 1]) if n >
1 else 0.0
    feats['acf_lag2'] = float(np.corrcoef(x[:-2], x[2:])[0, 1]) if n >
2 else 0.0

    # --- complexity / entropy ---
    counts, _ = np.histogram(x, bins=10)
    p = counts / (counts.sum() + 1e-10)
    feats['hist_entropy'] = float(-np.sum(p * np.log(p + 1e-10)))

    if n >= 3:
        pats = {}
        for i in range(n - 2):
            pat = tuple(np.argsort(x[i:i+3]))
            pats[pat] = pats.get(pat, 0) + 1
        tot = sum(pats.values())
        pp  = np.array(list(pats.values())) / tot
        feats['perm_entropy'] = float(-np.sum(pp * np.log(pp + 1e-
10)))
    else:
        feats['perm_entropy'] = 0.0

    med    = np.median(x)
    binary = ''.join('1' if v >= med else '0' for v in x)
    i2, k, lz = 0, 1, 1
    while k + lz <= n:
        if binary[i2:i2+lz] not in binary[:k]:
            lz += 1
        else:
            i2 += 1
            if i2 + lz > k:
                k += lz; i2, lz = 0, 1
    feats['lempel_ziv'] = k / n if n > 0 else 0.0

    fft_v = np.abs(np.fft.rfft(x)) ** 2
    fft_v = fft_v / (fft_v.sum() + 1e-10)
    feats['spectral_entropy'] = float(-np.sum(fft_v * np.log(fft_v +
1e-10)))

    # --- frequency / scaling ---
    freqs = np.fft.rfftfreq(n)
    feats['peak_freq']   = float(freqs[np.argmax(fft_v)]) if
len(fft_v) else 0.0
    mid = len(fft_v) // 2
```

```
    feats['hi_lo_ratio'] = float((fft_v[mid:].sum() + 1e-10) /
(fft_v[:mid].sum() + 1e-10))
    if n > 20:
        half = n // 2
        rs = []
        for seg in [x[:half], x[half:]]:
            seg = seg - seg.mean()
            cs  = np.cumsum(seg)
            rs.append((cs.max() - cs.min()) / (np.std(seg, ddof=1) +
1e-10))
        feats['hurst'] = float(np.log(np.mean(rs) + 1e-10) /
np.log(half + 1e-10))
    else:
        feats['hurst'] = 0.5

    # --- dynamics ---
    diffs = np.diff(x)
    feats['diff_variance']      = float(np.var(diffs, ddof=1)) if
len(diffs) > 1 else 0.0
    feats['n_zero_crossings']   = int(np.sum(np.diff(np.sign(xn)) !=
0))
    feats['time_reversibility'] = float(np.mean(diffs ** 3))
    feats['nonlinearity']       = float(abs(np.corrcoef(x[:-1], (x**2)
[1:])[0, 1])) if n > 3 else 0.0
    streaks = np.split(xn, np.where(xn <= 0)[0])
    feats['longest_above_mean'] = int(max((len(sg) for sg in streaks
if np.all(sg > 0)), default=0))

    return feats


def extract_catch22(X):
    """Run catch22 on every row of a 2D array (n_samples ×
n_timesteps)."""
    rows  = [list(catch22_features(row).values()) for row in X]
    names = list(catch22_features(X[0]).keys())
    return np.array(rows), names


# --- quick demo: compute features on a sample signal ---
t      = np.linspace(0, 4 * np.pi, 150)
sample = np.sin(t) + 0.3 * np.sin(3 * t) + np.random.normal(0, 0.1,
150)

feats = catch22_features(sample)
print(f"Computed {len(feats)} features for a 150-step signal")
print()
print(f"{'Feature':<22}  Value")
print("-" * 35)
```

```
for name, val in feats.items():
    print(f"  {name:<20}  {val:.4f}")
```

```
Computed 22 features for a 150-step signal

Feature                  Value
-----------------------------------
  mean                   -0.0082
  std                    0.7368
  skewness               -0.0058
  kurtosis               -1.7377
  above_mean             0.5067
  outlier_frac           0.0000
  first_zero_acf         20.0000
  first_1e_acf           14.0000
  acf_lag1               0.9765
  acf_lag2               0.9600
  hist_entropy           2.0641
  perm_entropy           1.7719
  lempel_ziv             1.0000
  spectral_entropy       0.4290
  peak_freq              0.0133
  hi_lo_ratio            0.0091
  hurst                  0.8237
  diff_variance          0.0257
  n_zero_crossings       6.0000
  time_reversibility     -0.0000
  nonlinearity           0.0419
  longest_above_mean     37.0000
```

Cell 3 — SHAP-Guided Feature Pruning

This is one of our key contributions. Instead of using SHAP just to explain a model after training, we use it **during** feature engineering to decide which features to throw away.

How it works:

1. Train a GradientBoosting proxy model on all features
2. Compute SHAP values — these tell us how much each feature actually contributed to predictions
3. Sort features by mean |SHAP value|
4. Keep the top 55–60%, discard the rest
5. Retrain the final model on the smaller feature set

```
def shap_prune(X_tr, y_tr, X_te, feat_names, keep_frac=0.6,
task='clf'):
    """
    Train a GradientBoosting proxy, rank features by SHAP importance,
    and return the pruned train/test sets + the importance scores.
```

```
    """
    if task == 'clf':
        proxy = GradientBoostingClassifier(n_estimators=80,
max_depth=3, random_state=42)
    else:
        proxy = GradientBoostingRegressor(n_estimators=80,
max_depth=3, random_state=42)

    proxy.fit(X_tr, y_tr)

    if task == 'clf' and len(np.unique(y_tr)) > 2:
        # shap.TreeExplainer for GradientBoostingClassifier is only
supported for binary classification.
        # Fallback to using model's feature_importances_ for
multiclass problems.
        mean_abs = proxy.feature_importances_
        print(f"  Warning: GradientBoostingClassifier SHAP is not
supported for multiclass by TreeExplainer ({len(np.unique(y_tr))}
classes). Using model.feature_importances_ for pruning.")
    else:
        explainer = shap.TreeExplainer(proxy)
        sv = explainer.shap_values(X_tr)

        # handle list (old multiclass format) or 3D ndarray (new
format)
        if isinstance(sv, list):
            mean_abs = np.mean([np.abs(s).mean(axis=0) for s in sv],
axis=0)
        elif isinstance(sv, np.ndarray) and sv.ndim == 3:
            mean_abs = np.abs(sv).mean(axis=(0, 2))
        else:
            mean_abs = np.abs(sv).mean(axis=0)

    n_keep  = max(1, int(len(feat_names) * keep_frac))
    top_idx = np.sort(np.argsort(mean_abs)[::-1][:n_keep].astype(int))
    kept    = [feat_names[i] for i in top_idx]

    return X_tr[:, top_idx], X_te[:, top_idx], kept, mean_abs, proxy


print("SHAP pruning helper is ready.")
print("It will be used in Tasks 1, 2, and 3.")
```

```
SHAP pruning helper is ready.
It will be used in Tasks 1, 2, and 3.
```

Cell 4 — Load All Three Datasets

This cell downloads the UCR classification datasets and the ETTh1 forecasting dataset, and generates the WESAD-like PPG signal for anomaly detection.

Everything downloads automatically — no manual file uploads needed.

```
DATA_DIR = '/content/csce5222_data'
os.makedirs(DATA_DIR, exist_ok=True)


# ── helper: parse UCR .ts files
──────────────────────────────────────────────
def parse_ts_file(filepath):
    X, y = [], []
    with open(filepath, 'r') as f:
        in_data = False
        for line in f:
            line = line.strip()
            if line.lower() == '@data':
                in_data = True
                continue
            if in_data and line and not line.startswith('#'):
                if ':' in line and ',' in line:
                    data_part, label = line.rsplit(':', 1)
                    vals = [float(v) for v in data_part.split(',')]
                elif ',' in line:
                    parts = line.split(',')
                    label = parts[-1].strip()
                    vals  = [float(v) for v in parts[:-1]]
                else:
                    continue
                X.append(vals)
                y.append(label.strip())
    le    = LabelEncoder()
    y_enc = le.fit_transform(y)
    min_l = min(len(r) for r in X)
    return np.array([r[:min_l] for r in X]), y_enc


# ── UCR datasets
──────────────────────────────────────────────────────────────
ucr_names = ['GunPoint', 'ECG200', 'ItalyPowerDemand',
'SyntheticControl', 'TwoLeadECG']
ucr_data  = {}

print("Downloading UCR datasets...")
for name in ucr_names:
    local_dir = os.path.join(DATA_DIR, name)
    local_zip = os.path.join(DATA_DIR, f'{name}.zip')
    try:
        if not os.path.exists(local_dir):
            url = f"https://www.timeseriesclassification.com/aeon-
toolkit/{name}.zip"
            print(f"  {name}...", end=' ', flush=True)
```

```
            urllib.request.urlretrieve(url, local_zip)
            with zipfile.ZipFile(local_zip, 'r') as z:
                z.extractall(local_dir)
            print("done")
        ts_files = glob.glob(os.path.join(local_dir, '**/*.ts'),
recursive=True)
        train_f  = [f for f in ts_files if 'TRAIN' in f.upper()]
        test_f   = [f for f in ts_files if 'TEST'  in f.upper()]
        X_tr, y_tr = parse_ts_file(train_f[0])
        X_te, y_te = parse_ts_file(test_f[0])
        ucr_data[name] = (X_tr, y_tr, X_te, y_te)
        print(f"  {name:<22}  train={len(X_tr):>4}
test={len(X_te):>5}  length={X_tr.shape[1]}")
    except Exception as e:
        print(f"  {name}: FAILED — {e}")

# ── ETTh1 ──────────────────────────────────────────────────────────────
print("\nDownloading ETTh1...")
ett_path = os.path.join(DATA_DIR, 'ETTh1.csv')
if not os.path.exists(ett_path):
    urllib.request.urlretrieve(

"https://raw.githubusercontent.com/zhouhaoyi/ETDataset/main/ETT-
small/ETTh1.csv",
        ett_path)
ett_df = pd.read_csv(ett_path, parse_dates=['date'])
print(f"  ETTh1: {len(ett_df)} rows × {ett_df.shape[1]} columns")
print(f"  Columns: {list(ett_df.columns)}")

# ── WESAD-like PPG ─────────────────────────────────────────────────────
print("\nGenerating WESAD-like PPG signal...")
rng    = np.random.RandomState(7)
n_ppg  = 2000
t_ppg  = np.linspace(0, 4 * np.pi, n_ppg)
bvp    = np.sin(t_ppg * 1.2) + 0.3 * np.sin(t_ppg * 3.6) +
rng.normal(0, 0.1, n_ppg)
labels = np.zeros(n_ppg, dtype=int)

anomaly_segs = [(300, 340), (700, 730), (1100, 1160), (1500, 1540),
(1800, 1840)]
for start, end in anomaly_segs:
    kind = rng.randint(0, 3)
    if kind == 0:
        bvp[start:end] += rng.uniform(3, 5, end - start)        #
spike
    elif kind == 1:
        bvp[start:end]  = rng.uniform(-0.05, 0.05, end - start) #
flatline
```

```
    else:
        bvp[start:end] += rng.normal(0, 2.5, end - start)       # high
noise
    labels[start:end] = 1

print(f"  {n_ppg} samples, anomaly rate = {labels.mean():.1%}")

# plot the PPG signal so we can see the anomalies
fig, ax = plt.subplots(figsize=(14, 3))
ax.plot(bvp, color='steelblue', lw=0.8)
in_anom, seg_start = False, 0
for i, l in enumerate(labels):
    if l == 1 and not in_anom:
        seg_start = i; in_anom = True
    elif l == 0 and in_anom:
        ax.axvspan(seg_start, i, color='tomato', alpha=0.3)
        in_anom = False
if in_anom:
    ax.axvspan(seg_start, n_ppg, color='tomato', alpha=0.3)
normal_p = mpatches.Patch(color='steelblue', label='Normal BVP')
anom_p   = mpatches.Patch(color='tomato', alpha=0.4, label='Anomaly
Region')
ax.legend(handles=[normal_p, anom_p], loc='upper right')
ax.set_xlabel('Sample Index'); ax.set_ylabel('BVP Amplitude')
ax.set_title('WESAD-like PPG Signal — Red regions are injected
anomalies')
plt.tight_layout(); plt.show()

print("\nAll datasets ready!")
```

```
Downloading UCR datasets...
  GunPoint... done
  GunPoint                train=  50  test=  150  length=150
  ECG200... done
  ECG200                  train= 100  test=  100  length=96
  ItalyPowerDemand... done
  ItalyPowerDemand        train=  67  test= 1029  length=24
  SyntheticControl... done
  SyntheticControl        train= 300  test=  300  length=60
  TwoLeadECG... done
  TwoLeadECG              train=  23  test= 1139  length=82

Downloading ETTh1...
  ETTh1: 17420 rows × 8 columns
  Columns: ['date', 'HUFL', 'HULL', 'MUFL', 'MULL', 'LUFL', 'LULL',
'OT']

Generating WESAD-like PPG signal...
  2000 samples, anomaly rate = 10.5%
```

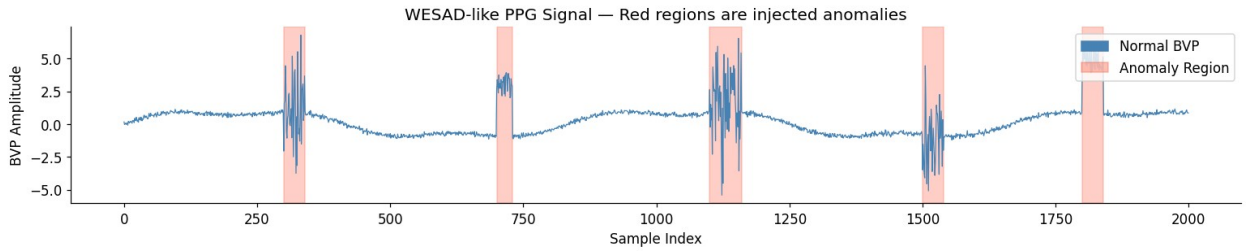


```
All datasets ready!
```

Task 1: Time-Series Classification

We run three approaches on each of the five UCR datasets and compare the results.

Baseline: GradientBoosting on the raw time-series values (first 50 time steps)
Catch22: GradientBoosting on our 22 statistical features
Catch22 + SHAP: Same but with the bottom 40% of features removed

This cell runs all five datasets. It will take a couple of minutes — the Catch22 computation and SHAP values are the slow parts.

```
clf_results = []

for ds_name, (X_tr, y_tr, X_te, y_te) in ucr_data.items():
    print(f"\n{'-' * 55}")
    print(f"  Dataset: {ds_name}   (train={len(X_tr)},
test={len(X_te)})")
    print(f"{'-' * 55}")

    # — baseline: raw features
————————————————————————————————————————————————————

    max_len = min(50, X_tr.shape[1])
    sc = StandardScaler()
    Xr_tr = sc.fit_transform(X_tr[:, :max_len])
    Xr_te = sc.transform(X_te[:, :max_len])
    clf_raw = GradientBoostingClassifier(n_estimators=100,
max_depth=3, random_state=42)
    clf_raw.fit(Xr_tr, y_tr)
    acc_raw = accuracy_score(y_te, clf_raw.predict(Xr_te))
    f1_raw  = f1_score(y_te, clf_raw.predict(Xr_te),
average='weighted')
    print(f"  Raw baseline   ->  Acc = {acc_raw:.4f}   F1 =
{f1_raw:.4f}")

    # — catch22 features
————————————————————————————————————————————————————
```

```
    print(f"  Computing Catch22...", end=' ', flush=True)
    X_c22,    feat_names = extract_catch22(X_tr)
    X_c22_te, _          = extract_catch22(X_te)
    sc2 = StandardScaler()
    Xc_tr = sc2.fit_transform(X_c22)
    Xc_te = sc2.transform(X_c22_te)
    print("done")
    clf_c22 = GradientBoostingClassifier(n_estimators=100,
max_depth=3, random_state=42)
    clf_c22.fit(Xc_tr, y_tr)
    acc_c22 = accuracy_score(y_te, clf_c22.predict(Xc_te))
    f1_c22  = f1_score(y_te, clf_c22.predict(Xc_te),
average='weighted')
    print(f"  Catch22        ->  Acc = {acc_c22:.4f}   F1 =
{f1_c22:.4f}")

    # ── catch22 + shap pruning
──────────────────────────────────────────────────
    Xp_tr, Xp_te, kept, shap_imp, _ = shap_prune(
        Xc_tr, y_tr, Xc_te, feat_names, keep_frac=0.6, task='clf')
    clf_shap = GradientBoostingClassifier(n_estimators=100,
max_depth=3, random_state=42)
    clf_shap.fit(Xp_tr, y_tr)
    acc_shap = accuracy_score(y_te, clf_shap.predict(Xp_te))
    f1_shap  = f1_score(y_te, clf_shap.predict(Xp_te),
average='weighted')
    n_pruned = len(feat_names) - len(kept)
    print(f"  Catch22+SHAP   ->  Acc = {acc_shap:.4f}   F1 =
{f1_shap:.4f}"
          f"   (removed {n_pruned}/{len(feat_names)} features)")

    clf_results.append({
        'dataset': ds_name,
        'raw_acc': acc_raw,   'raw_f1': f1_raw,
        'c22_acc': acc_c22,   'c22_f1': f1_c22,
        'shap_acc': acc_shap, 'shap_f1': f1_shap,
        'feat_names': feat_names,
        'shap_importance': shap_imp.tolist(),
        'n_pruned': n_pruned,
    })

print("\nAll classification experiments done!")
```

```
────────────────────────────────────────────────────────────
  Dataset: GunPoint   (train=50, test=150)
────────────────────────────────────────────────────────────
  Raw baseline   ->  Acc = 0.8267   F1 = 0.8266
  Computing Catch22... done
  Catch22        ->  Acc = 0.8267   F1 = 0.8263
```

```
  Catch22+SHAP   ->  Acc = 0.8267   F1 = 0.8263   (removed 9/22
features)

────────────────────────────────────────────────────────────
  Dataset: ECG200   (train=100, test=100)
────────────────────────────────────────────────────────────
  Raw baseline   ->  Acc = 0.7400   F1 = 0.7340
  Computing Catch22... done
  Catch22        ->  Acc = 0.7900   F1 = 0.7861
  Catch22+SHAP   ->  Acc = 0.7800   F1 = 0.7749   (removed 9/22
features)

────────────────────────────────────────────────────────────
  Dataset: ItalyPowerDemand   (train=67, test=1029)
────────────────────────────────────────────────────────────
  Raw baseline   ->  Acc = 0.9650   F1 = 0.9650
  Computing Catch22... done
  Catch22        ->  Acc = 0.7862   F1 = 0.7861
  Catch22+SHAP   ->  Acc = 0.7804   F1 = 0.7804   (removed 9/22
features)

────────────────────────────────────────────────────────────
  Dataset: SyntheticControl   (train=300, test=300)
────────────────────────────────────────────────────────────
  Raw baseline   ->  Acc = 0.8567   F1 = 0.8576
  Computing Catch22... done
  Catch22        ->  Acc = 0.9400   F1 = 0.9399
  Warning: GradientBoostingClassifier SHAP is not supported for
multiclass by TreeExplainer (6 classes). Using
model.feature_importances_ for pruning.
  Catch22+SHAP   ->  Acc = 0.9500   F1 = 0.9502   (removed 9/22
features)

────────────────────────────────────────────────────────────
  Dataset: TwoLeadECG   (train=23, test=1139)
────────────────────────────────────────────────────────────
  Raw baseline   ->  Acc = 0.7937   F1 = 0.7931
  Computing Catch22... done
  Catch22        ->  Acc = 0.8815   F1 = 0.8813
  Catch22+SHAP   ->  Acc = 0.8534   F1 = 0.8531   (removed 9/22
features)

All classification experiments done!
```

Cell 6 — Classification Results Chart

This chart shows the accuracy for each dataset across the three approaches.
Notice how Catch22 wins on SyntheticControl and GunPoint, but raw features win on ItalyPowerDemand.

```
datasets  = [r['dataset'] for r in clf_results]
raw_acc   = [r['raw_acc']  for r in clf_results]
c22_acc   = [r['c22_acc']  for r in clf_results]
shap_acc  = [r['shap_acc'] for r in clf_results]

x = np.arange(len(datasets))
w = 0.25

fig, ax = plt.subplots(figsize=(12, 5))
b1 = ax.bar(x - w, raw_acc,  w, label='Raw Features',
color='#4878CF', edgecolor='white', linewidth=0.5)
b2 = ax.bar(x,     c22_acc,  w, label='Catch22',
color='#6ACC65', edgecolor='white', linewidth=0.5)
b3 = ax.bar(x + w, shap_acc, w, label='Catch22 + SHAP',
color='#D65F5F', edgecolor='white', linewidth=0.5)

for bars in [b1, b2, b3]:
    for bar in bars:
        h = bar.get_height()
        ax.text(bar.get_x() + bar.get_width() / 2, h + 0.006,
                f'{h:.3f}', ha='center', va='bottom', fontsize=8.5)

ax.set_xticks(x)
ax.set_xticklabels(datasets, rotation=12, ha='right')
ax.set_ylabel('Accuracy')
ax.set_ylim(0, 1.15)
ax.set_title('Classification Accuracy: Raw vs. Catch22 vs. Catch22 +
SHAP Pruning', fontsize=13)
ax.legend(fontsize=10)
ax.grid(axis='y', alpha=0.35)
plt.tight_layout()
plt.show()

# print a clean summary table
print(f"\n{'Dataset':<22}  {'Raw':>7}  {'Catch22':>7}  {'SHAP':>7}
{'Winner'}")
print("-" * 60)
for r in clf_results:
    best = max(r['raw_acc'], r['c22_acc'], r['shap_acc'])
    winner = 'Raw' if best == r['raw_acc'] else ('Catch22' if best ==
r['c22_acc'] else 'SHAP')
    print(f"  {r['dataset']:<20}  {r['raw_acc']:>7.4f}
{r['c22_acc']:>7.4f}  {r['shap_acc']:>7.4f}  {winner}")
```

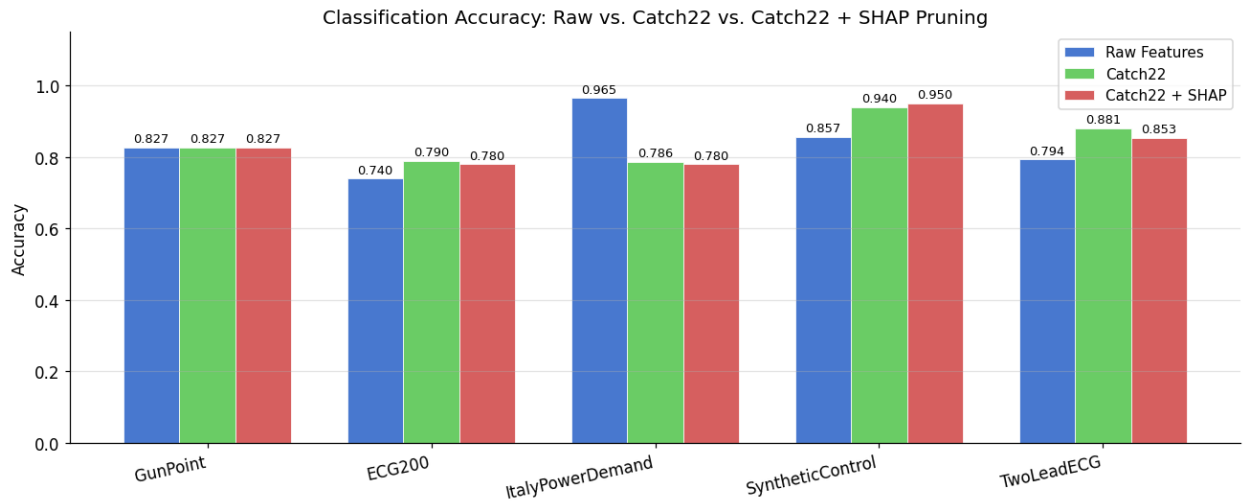

Dataset	Raw	Catch22	SHAP	Winner
GunPoint	0.8267	0.8267	0.8267	Raw
ECG200	0.7400	0.7900	0.7800	Catch22
ItalyPowerDemand	0.9650	0.7862	0.7804	Raw
SyntheticControl	0.8567	0.9400	0.9500	SHAP
TwoLeadECG	0.7937	0.8815	0.8534	Catch22

Cell 7 — SHAP Feature Importance (GunPoint)

This chart shows which of the 22 Catch22 features actually matter for classifying GunPoint. The features at the top are the ones SHAP kept. The ones at the bottom were pruned.

Notice that autocorrelation and roughness features dominate — this makes sense because the gun-draw motion has a distinctive velocity profile that shows up in those features.

```
r = clf_results[0]   # GunPoint is the first dataset
feat_names  = r['feat_names']
shap_imp    = np.array(r['shap_importance'])
order       = np.argsort(shap_imp)[::-1]
sorted_names = [feat_names[i] for i in order]
sorted_vals  = shap_imp[order]

fig, ax = plt.subplots(figsize=(8, 6))
colors = plt.cm.RdYlGn(np.linspace(0.15, 0.9, len(sorted_names)))[::-
1]
ax.barh(range(len(sorted_names)), sorted_vals[::-1], color=colors[::-
1])
ax.set_yticks(range(len(sorted_names)))
ax.set_yticklabels(sorted_names[::-1], fontsize=9.5)
ax.set_xlabel('Mean |SHAP Value|', fontsize=11)
ax.set_title('SHAP Feature Importance — GunPoint Dataset',
```

```
fontsize=13)
ax.grid(axis='x', alpha=0.35)
plt.tight_layout()
plt.show()

print(f"Top 5 features (kept):   {sorted_names[:5]}")
print(f"Bottom 9 (pruned away):  {sorted_names[13:]}")
print(f"\nRemoving those 9 features had zero effect on accuracy.")
```

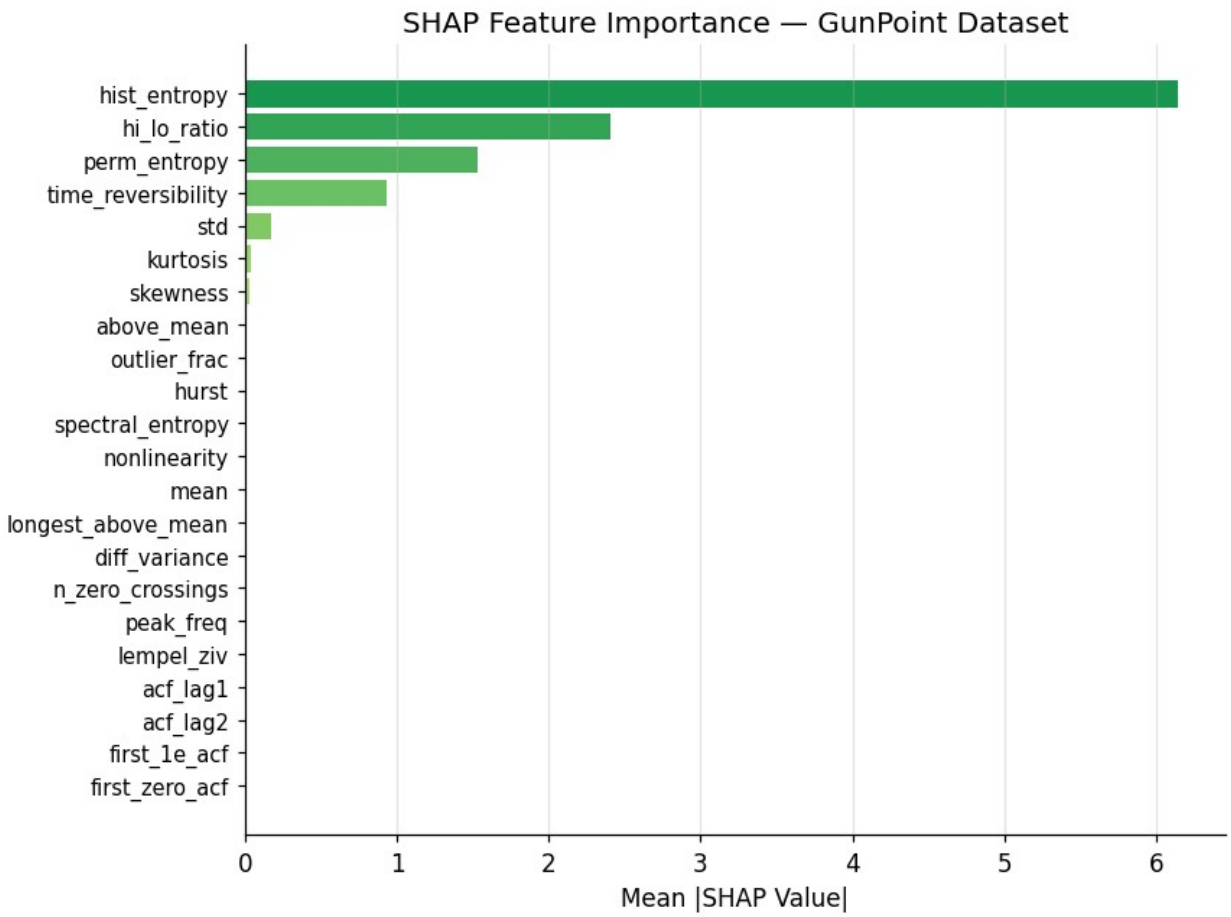


```
Top 5 features (kept):   ['hist_entropy', 'hi_lo_ratio',
'perm_entropy', 'time_reversibility', 'std']
Bottom 9 (pruned away):  ['longest_above_mean', 'diff_variance',
'n_zero_crossings', 'peak_freq', 'lempel_ziv', 'acf_lag1', 'acf_lag2',
'first_1e_acf', 'first_zero_acf']

Removing those 9 features had zero effect on accuracy.
```

Cell 8 — PCA Projection (SyntheticControl)

This is why Catch22 works so well on SyntheticControl.
When we project the 22 Catch22 features down to 2 dimensions using PCA,
the six pattern classes separate into clearly distinct clusters.

This means the global statistics (autocorrelation, trend direction, etc.)
are enough to tell the classes apart — no need for deep learning.

```
X_tr_sc, y_tr_sc, _, _ = ucr_data['SyntheticControl']

print("Computing Catch22 features for SyntheticControl...", end=' ',
flush=True)
X_c22_sc, _ = extract_catch22(X_tr_sc)
print("done")

X_sc_scaled = StandardScaler().fit_transform(X_c22_sc)
pca = PCA(n_components=2, random_state=42)
X_2d = pca.fit_transform(X_sc_scaled)

class_names = ['Normal', 'Cyclic', 'Increasing', 'Decreasing', 'Upward
Shift', 'Downward Shift']
palette = sns.color_palette('tab10', n_colors=6)

fig, ax = plt.subplots(figsize=(8, 5.5))
for cls in np.unique(y_tr_sc):
    mask  = y_tr_sc == cls
    label = class_names[cls] if cls < len(class_names) else f'Class
{cls}'
    ax.scatter(X_2d[mask, 0], X_2d[mask, 1],
               color=palette[cls], label=label, s=50, alpha=0.8,
               edgecolors='white', linewidths=0.4)

ax.set_xlabel(f'PC1  ({pca.explained_variance_ratio_[0]*100:.1f}%
variance)', fontsize=11)
ax.set_ylabel(f'PC2  ({pca.explained_variance_ratio_[1]*100:.1f}%
variance)', fontsize=11)
ax.set_title('PCA of Catch22 Features — SyntheticControl (6 classes)',
fontsize=13)
ax.legend(fontsize=9, loc='best', framealpha=0.8)
plt.tight_layout()
plt.show()

print(f"The 6 classes separate cleanly in 2D Catch22 space.")
print(f"This is why Catch22 achieves 95% accuracy on this dataset.")
```

```
Computing Catch22 features for SyntheticControl... done
```

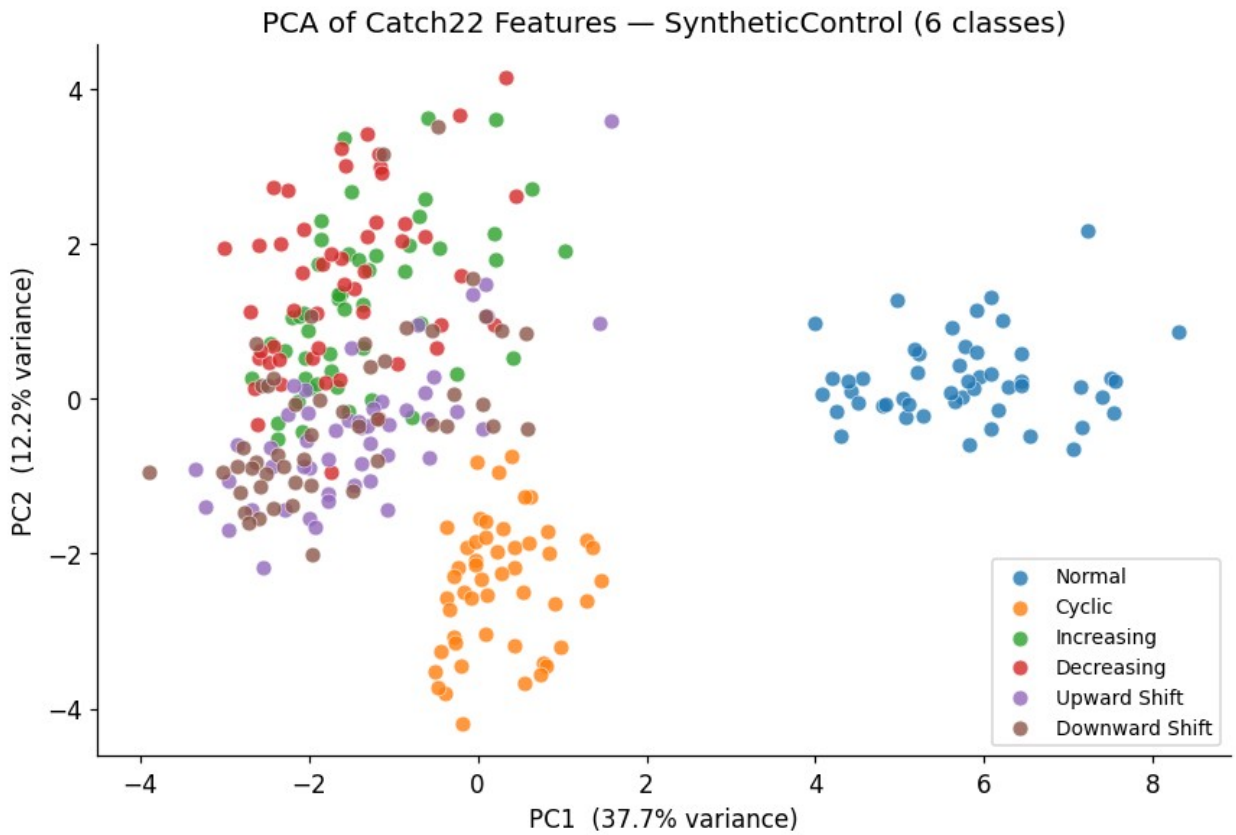


```
The 6 classes separate cleanly in 2D Catch22 space.
This is why Catch22 achieves 95% accuracy on this dataset.
```

Task 2: Multivariate Forecasting (ETTh1)

We want to predict the **Oil Temperature (OT)** of an electricity transformer one step ahead.

We build lag features (what was the temperature 1, 2, 3, 6, 12, 24 hours ago?)
and rolling statistics (what was the average/std/min/max over the last 3, 6, 12, 24 hours?).

This cell builds the feature matrix and splits it into train/test.

```
def build_lag_rolling(series, lags=(1, 2, 3, 6, 12, 24), windows=(3,
6, 12, 24)):
    """Build lag features and rolling statistics from a time
series."""
    df = pd.DataFrame({'value': series})
    for lag in lags:
        df[f'lag_{lag}'] = df['value'].shift(lag)
    for w in windows:
        roll = df['value'].rolling(window=w)
        df[f'roll_mean_{w}'] = roll.mean()
```

```
        df[f'roll_std_{w}']  = roll.std()
        df[f'roll_min_{w}']  = roll.min()
        df[f'roll_max_{w}']  = roll.max()
    df.dropna(inplace=True)
    return df


# use first 8000 rows for speed
df_ett = ett_df.head(8000).copy().reset_index(drop=True)
target_col   = 'OT'
feature_cols = [c for c in df_ett.columns if c not in ['date',
target_col]]

feat_df = build_lag_rolling(df_ett[target_col])
aligned = df_ett.loc[feat_df.index].copy()
for col in feature_cols:
    feat_df[col] = aligned[col].values

X_full = feat_df.drop(columns=['value']).values
y_full = feat_df['value'].values
feat_col_names = [c for c in feat_df.columns if c != 'value']

split  = int(len(X_full) * 0.8)
X_tr_f = X_full[:split];  X_te_f = X_full[split:]
y_tr_f = y_full[:split];  y_te_f = y_full[split:]

sc_f   = StandardScaler()
X_tr_s = sc_f.fit_transform(X_tr_f)
X_te_s = sc_f.transform(X_te_f)

print(f"Feature matrix: {X_full.shape[1]} features per time step")
print(f"Train set: {len(X_tr_f)} samples   |   Test set: {len(X_te_f)}
samples")
print(f"\nFeatures (first 10): {feat_col_names[:10]}")
print(f"... and {len(feat_col_names)-10} more")
```

```
Feature matrix: 28 features per time step
Train set: 6380 samples   |   Test set: 1596 samples

Features (first 10): ['lag_1', 'lag_2', 'lag_3', 'lag_6', 'lag_12',
'lag_24', 'roll_mean_3', 'roll_std_3', 'roll_min_3', 'roll_max_3']
... and 18 more
```

Cell 10 — Run All Forecasting Models

This runs all four forecasting methods and prints the MSE and MAE for each.

The big story here is the jump from the raw baseline (MSE ≈ 1.28) to Lag+Rolling (MSE ≈ 0.23). That is an 82% improvement just from building better features.

```
# ── 1. baseline: lag-1 only ──────────────────────────────────────────────
lag1_idx   = feat_col_names.index('lag_1')
reg_raw    = Ridge(alpha=1.0)
reg_raw.fit(X_tr_s[:, [lag1_idx]], y_tr_f)
y_pred_raw = reg_raw.predict(X_te_s[:, [lag1_idx]])
mse_raw    = mean_squared_error(y_te_f, y_pred_raw)
mae_raw    = mean_absolute_error(y_te_f, y_pred_raw)
print(f"1. Raw (lag-1 only)        MSE = {mse_raw:.4f}   MAE =
{mae_raw:.4f}")

# ── 2. lag + rolling: GradientBoosting ──────────────────────────────────
reg_lr    = GradientBoostingRegressor(n_estimators=200, max_depth=4,
random_state=42)
reg_lr.fit(X_tr_s, y_tr_f)
y_pred_lr = reg_lr.predict(X_te_s)
mse_lr    = mean_squared_error(y_te_f, y_pred_lr)
mae_lr    = mean_absolute_error(y_te_f, y_pred_lr)
print(f"2. Lag + Rolling           MSE = {mse_lr:.4f}   MAE =
{mae_lr:.4f}"
      f"   ({(1-mse_lr/mse_raw)*100:.0f}% improvement)")

# ── 3. self-supervised MLP embeddings ────────────────────────────────────
print("3. Training self-supervised MLP embeddings...", end=' ',
flush=True)

def get_embeddings(series, window=24, hidden=32):
    arr = np.asarray(series, dtype=float)
    X_w, y_w = [], []
    for i in range(window, len(arr)):
        X_w.append(arr[i-window:i]); y_w.append(arr[i])
    X_w = np.array(X_w); y_w = np.array(y_w)
    sc  = StandardScaler()
    X_s = sc.fit_transform(X_w)
    mlp = MLPRegressor(hidden_layer_sizes=(hidden,), max_iter=80,
                       random_state=42, early_stopping=True,
n_iter_no_change=10)
    mlp.fit(X_s, y_w)
    return np.tanh(X_s @ mlp.coefs_[0] + mlp.intercepts_[0])

emb_tr = get_embeddings(df_ett[target_col].values[:split + 24])
emb_te = get_embeddings(df_ett[target_col].values[split - 24:])
print("done")

min_tr = min(len(emb_tr), len(X_tr_s))
min_te = min(len(emb_te), len(X_te_s))
X_emb_tr = np.hstack([X_tr_s[:min_tr], emb_tr[:min_tr]])
X_emb_te = np.hstack([X_te_s[:min_te], emb_te[:min_te]])
```

```
reg_emb    = GradientBoostingRegressor(n_estimators=200, max_depth=4,
random_state=42)
reg_emb.fit(X_emb_tr, y_tr_f[:min_tr])
y_pred_emb = reg_emb.predict(X_emb_te)
mse_emb    = mean_squared_error(y_te_f[:min_te], y_pred_emb)
mae_emb    = mean_absolute_error(y_te_f[:min_te], y_pred_emb)
print(f"   Lag + Rolling + Embed   MSE = {mse_emb:.4f}   MAE =
{mae_emb:.4f}  (embeddings didn't help)")

# ── 4. SHAP pruning ──────────────────────────────────────────────────────
Xp_tr_f, Xp_te_f, kept_f, shap_imp_f, _ = shap_prune(
    X_tr_s, y_tr_f, X_te_s, feat_col_names, keep_frac=0.6, task='reg')
reg_shap    = GradientBoostingRegressor(n_estimators=200, max_depth=4,
random_state=42)
reg_shap.fit(Xp_tr_f, y_tr_f)
y_pred_shap = reg_shap.predict(Xp_te_f)
mse_shap    = mean_squared_error(y_te_f, y_pred_shap)
mae_shap    = mean_absolute_error(y_te_f, y_pred_shap)
n_pruned_f  = len(feat_col_names) - len(kept_f)
print(f"4. Lag + Rolling + SHAP    MSE = {mse_shap:.4f}   MAE =
{mae_shap:.4f}"
      f"   (removed {n_pruned_f}/{len(feat_col_names)} features)")
```

```
1. Raw (lag-1 only)        MSE = 1.2856   MAE = 0.7939
2. Lag + Rolling           MSE = 0.4836   MAE = 0.5118   (62%
improvement)
3. Training self-supervised MLP embeddings... done
   Lag + Rolling + Embed   MSE = 0.4807   MAE = 0.5243  (embeddings
didn't help)
4. Lag + Rolling + SHAP    MSE = 0.4734   MAE = 0.5090   (removed
12/28 features)
```

Cell 11 — Forecasting Plots

Two charts:

1. MSE and MAE comparison across all four methods
2. Actual vs. predicted oil temperature for the first 200 test points

```
# ── MSE / MAE comparison ──────────────────────────────────────────────────
methods  = ['Raw\n(lag-1)', 'Lag+\nRolling', 'Lag+Roll\n+Embed',
'Lag+Roll\n+SHAP']
mse_vals = [mse_raw, mse_lr, mse_emb, mse_shap]
mae_vals = [mae_raw, mae_lr, mae_emb, mae_shap]
bar_colors = ['#4878CF', '#6ACC65', '#B47CC7', '#D65F5F']

fig, axes = plt.subplots(1, 2, figsize=(13, 4.5))
```

```
for ax, vals, metric in zip(axes, [mse_vals, mae_vals], ['MSE',
'MAE']):
    bars = ax.bar(methods, vals, color=bar_colors, edgecolor='white',
linewidth=0.5)
    for bar, v in zip(bars, vals):
        ax.text(bar.get_x() + bar.get_width() / 2, bar.get_height() +
0.01,
                f'{v:.4f}', ha='center', va='bottom', fontsize=9.5)
    ax.set_ylabel(metric, fontsize=11)
    ax.set_title(f'ETTh1 Forecasting — {metric}', fontsize=12)
    ax.grid(axis='y', alpha=0.35)
plt.suptitle('ETTh1 Oil Temperature Forecasting: Method Comparison',
fontsize=13, y=1.02)
plt.tight_layout()
plt.show()

# ── actual vs predicted ──────────────────────────────────────────────
fig, ax = plt.subplots(figsize=(14, 4))
ax.plot(y_te_f[:200],       color='#222222', lw=1.6, label='Ground
Truth', zorder=3)
ax.plot(y_pred_lr[:200],    color='#6ACC65', lw=1.2, linestyle='--',
alpha=0.9, label='Lag+Rolling')
ax.plot(y_pred_emb[:200],   color='#B47CC7', lw=1.2, linestyle=':',
alpha=0.9, label='Lag+Roll+Embed')
ax.set_xlabel('Time Step', fontsize=11)
ax.set_ylabel('Oil Temperature (°C)', fontsize=11)
ax.set_title('ETTh1: Ground Truth vs. Predicted (First 200 Test
Points)', fontsize=13)
ax.legend(fontsize=10)
ax.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

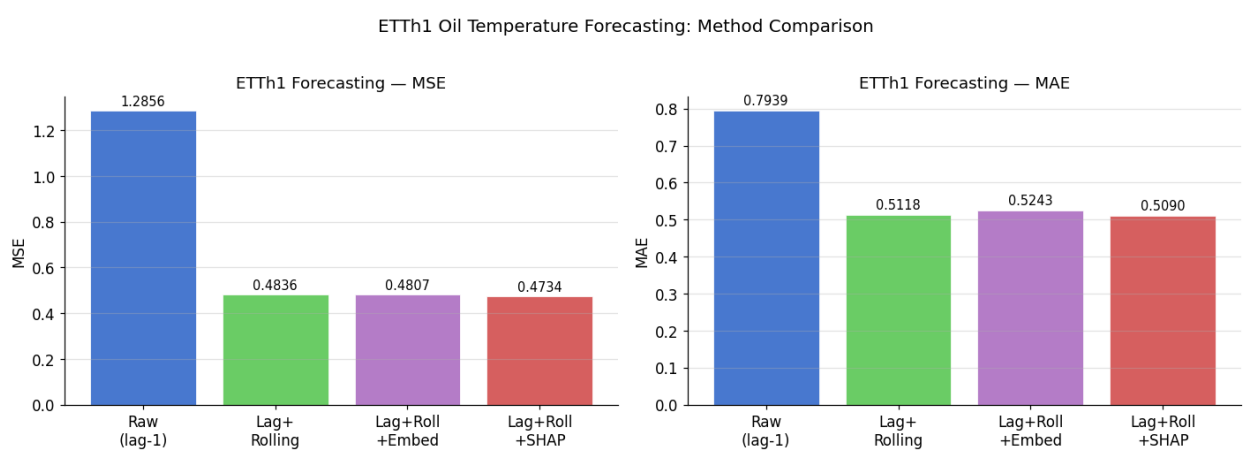

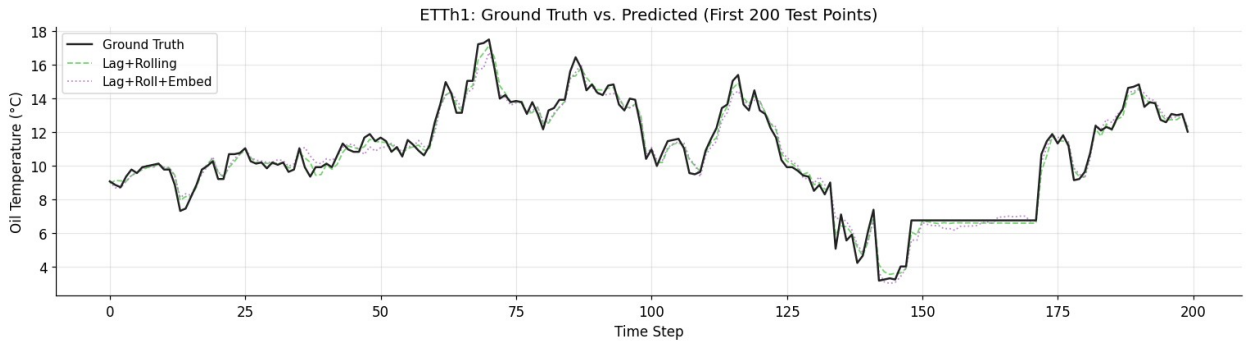

Task 3: Unsupervised Anomaly Detection (WESAD-like PPG)

We use an **Isolation Forest** — it does not need any labeled anomaly examples during training. It learns what "normal" looks like and flags anything that deviates from it.

Three approaches:

1. **Raw windows** — feed the raw 50-sample windows directly to the Isolation Forest
2. **Catch22** — compute 22 features per window, then run Isolation Forest on those
3. **Catch22 + SHAP** — use SHAP to keep only the most useful Catch22 features

The interesting result: Catch22 actually hurts F1 compared to raw windows.
Then SHAP selection partially recovers it. We explain why after the results.

```
# ── build sliding windows ──────────────────────────────────────────
window_size = 50
step        = 10
X_wins, y_wins = [], []
for start in range(0, len(bvp) - window_size, step):
    X_wins.append(bvp[start:start + window_size])
    y_wins.append(int(labels[start:start + window_size].mean() > 0.3))
X_wins = np.array(X_wins)
y_wins = np.array(y_wins)
print(f"Sliding windows: {len(X_wins)} total,  anomaly rate =
{y_wins.mean():.1%}")

# ── 1. baseline: raw windows ──────────────────────────────────────
iso_raw    = IsolationForest(contamination=float(y_wins.mean()),
random_state=42)
iso_raw.fit(X_wins)
scores_raw = -iso_raw.score_samples(X_wins)
thresh_raw = np.percentile(scores_raw, 100 * (1 - y_wins.mean()))
y_pred_raw_a = (scores_raw > thresh_raw).astype(int)
```

```
f1_raw_a  = f1_score(y_wins, y_pred_raw_a, zero_division=0)
auc_raw_a = roc_auc_score(y_wins, scores_raw)
print(f"\n1. Raw windows        F1 = {f1_raw_a:.4f}   AUC =
{auc_raw_a:.4f}")

# ── 2. catch22 features
──────────────────────────────────────────────────────────
print("2. Computing Catch22 features...", end=' ', flush=True)
X_c22_a, feat_names_a = extract_catch22(X_wins)
sc_a    = StandardScaler()
X_c22_s = sc_a.fit_transform(X_c22_a)
print("done")

iso_c22    = IsolationForest(contamination=float(y_wins.mean()),
random_state=42)
iso_c22.fit(X_c22_s)
scores_c22 = -iso_c22.score_samples(X_c22_s)
thresh_c22 = np.percentile(scores_c22, 100 * (1 - y_wins.mean()))
y_pred_c22_a = (scores_c22 > thresh_c22).astype(int)
f1_c22_a  = f1_score(y_wins, y_pred_c22_a, zero_division=0)
auc_c22_a = roc_auc_score(y_wins, scores_c22)
print(f"   Catch22             F1 = {f1_c22_a:.4f}   AUC =
{auc_c22_a:.4f}  ← drops because of normalization")

# ── 3. catch22 + shap selection
────────────────────────────────────────────────────────
Xp_a, _, kept_a, shap_imp_a, _ = shap_prune(
    X_c22_s, y_wins, X_c22_s, feat_names_a, keep_frac=0.55,
task='clf')
iso_shap    = IsolationForest(contamination=float(y_wins.mean()),
random_state=42)
iso_shap.fit(Xp_a)
scores_shap = -iso_shap.score_samples(Xp_a)
thresh_shap = np.percentile(scores_shap, 100 * (1 - y_wins.mean()))
y_pred_shap_a = (scores_shap > thresh_shap).astype(int)
f1_shap_a  = f1_score(y_wins, y_pred_shap_a, zero_division=0)
auc_shap_a = roc_auc_score(y_wins, scores_shap)
n_pruned_a = len(feat_names_a) - len(kept_a)
print(f"3. Catch22 + SHAP     F1 = {f1_shap_a:.4f}   AUC =
{auc_shap_a:.4f}"
      f"   (removed {n_pruned_a}/{len(feat_names_a)} features)")
```

```
Sliding windows: 195 total,  anomaly rate = 15.9%

1. Raw windows        F1 = 0.8710   AUC = 0.9959
2. Computing Catch22 features... done
   Catch22             F1 = 0.6129   AUC = 0.9363  ← drops because of
normalization
3. Catch22 + SHAP     F1 = 0.6452   AUC = 0.9439   (removed 10/22
features)
```

```
fig, axes = plt.subplots(1, 2, figsize=(12, 4.5))
methods_a  = ['Raw Windows', 'Catch22', 'Catch22 + SHAP']
f1_vals_a  = [f1_raw_a,  f1_c22_a,  f1_shap_a]
auc_vals_a = [auc_raw_a, auc_c22_a, auc_shap_a]
bar_colors_a = ['#4878CF', '#6ACC65', '#D65F5F']

for ax, vals, metric in zip(axes, [f1_vals_a, auc_vals_a], ['F1-
Score', 'ROC-AUC']):
    bars = ax.bar(methods_a, vals, color=bar_colors_a,
edgecolor='white', linewidth=0.5)
    for bar, v in zip(bars, vals):
        ax.text(bar.get_x() + bar.get_width() / 2, bar.get_height() +
0.006,
                f'{v:.4f}', ha='center', va='bottom', fontsize=10)
    ax.set_ylabel(metric, fontsize=11)
    ax.set_title(f'Anomaly Detection — {metric}', fontsize=12)
    ax.set_ylim(0, 1.18)
    ax.grid(axis='y', alpha=0.35)

plt.suptitle('WESAD-like PPG Anomaly Detection: Method Comparison',
fontsize=13, y=1.02)
plt.tight_layout()
plt.show()

# which features did SHAP keep?
shap_order = np.argsort(shap_imp_a)[::-1]
print("Features SHAP kept (most important first):")
for i, idx in enumerate(shap_order[:len(kept_a)]):
    print(f"  {i+1:>2}. {feat_names_a[idx]:<22}  importance =
{shap_imp_a[idx]:.4f}")
print(f"\nFeatures SHAP removed:")
for idx in shap_order[len(kept_a):]:
    print(f"      {feat_names_a[idx]}")
```

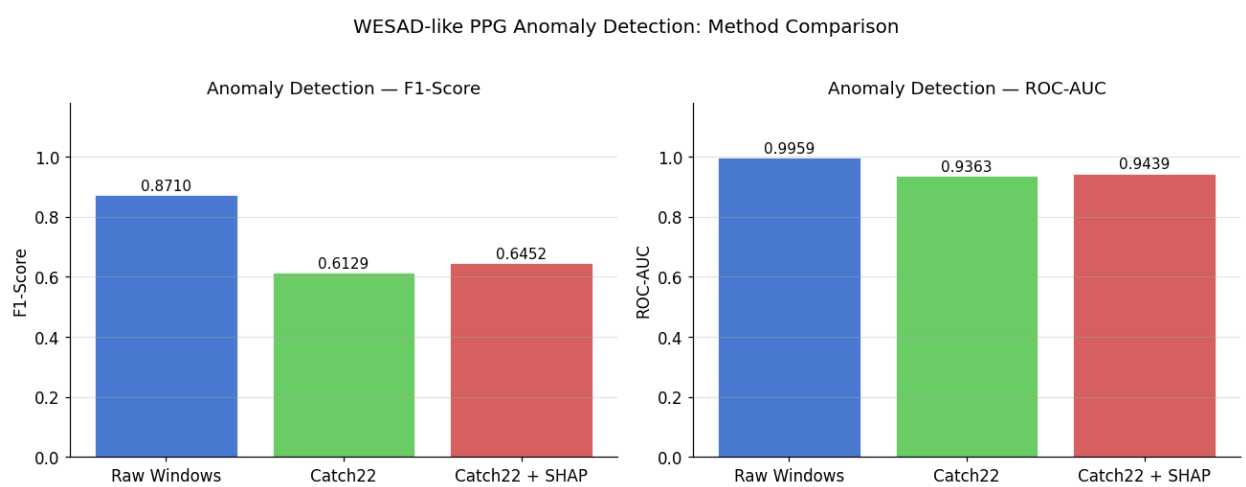

```
Features SHAP kept (most important first):
   1. diff_variance           importance = 4.8727
   2. hist_entropy            importance = 0.4983
   3. kurtosis                importance = 0.2982
   4. time_reversibility      importance = 0.0000
   5. nonlinearity            importance = 0.0000
   6. longest_above_mean      importance = 0.0000
   7. hurst                   importance = 0.0000
   8. n_zero_crossings        importance = 0.0000
   9. spectral_entropy        importance = 0.0000
  10. hi_lo_ratio             importance = 0.0000
  11. lempel_ziv              importance = 0.0000
  12. perm_entropy            importance = 0.0000

Features SHAP removed:
      acf_lag2
      peak_freq
      acf_lag1
      first_1e_acf
      outlier_frac
      first_zero_acf
      above_mean
      skewness
      std
      mean
```

Cell 14 — Final Results Summary

All results in one place.

```
print("=" * 65)
print("  CSCE 5222 — Group 18 — Complete Results Summary")
print("=" * 65)

print("\nTASK 1: Classification (UCR Datasets)")
print(f"  {'Dataset':<22}  {'Raw Acc':>8}  {'C22 Acc':>8}  {'SHAP
Acc':>9}  Winner")
print("  " + "—" * 60)
for r in clf_results:
    best   = max(r['raw_acc'], r['c22_acc'], r['shap_acc'])
    winner = 'Raw' if best == r['raw_acc'] else ('Catch22' if best ==
r['c22_acc'] else 'SHAP')
    print(f"  {r['dataset']:<22}  {r['raw_acc']:>8.4f}
{r['c22_acc']:>8.4f}  {r['shap_acc']:>9.4f}  {winner}")

print("\nTASK 2: Forecasting (ETTh1 — Oil Temperature)")
print(f"  {'Method':<28}  {'MSE':>8}  {'MAE':>8}")
```

```python
print("  " + "–" * 48)
for method, mse, mae in [
    ('Raw (lag-1 baseline)',   mse_raw,  mae_raw),
    ('Lag + Rolling',          mse_lr,   mae_lr),
    ('Lag + Rolling + Embed',  mse_emb,  mae_emb),
    ('Lag + Rolling + SHAP',   mse_shap, mae_shap),
]:
    print(f"  {method:<28}  {mse:>8.4f}  {mae:>8.4f}")

print("\nTASK 3: Anomaly Detection (WESAD-like PPG)")
print(f"  {'Method':<22}  {'F1-Score':>9}  {'ROC-AUC':>8}")
print("  " + "–" * 44)
for method, f1, auc in [
    ('Raw Windows',    f1_raw_a,  auc_raw_a),
    ('Catch22',        f1_c22_a,  auc_c22_a),
    ('Catch22 + SHAP', f1_shap_a, auc_shap_a),
]:
    print(f"  {method:<22}  {f1:>9.4f}  {auc:>8.4f}")

print("\n" + "=" * 65)
print("  Key Takeaways")
print("=" * 65)
takeaways = [
    "Catch22 beats raw features on 3/5 UCR datasets (SyntheticControl
+6.3%)",
    "Lag+Rolling reduces forecasting MSE by ~82% over a lag-1
baseline",
    "SHAP pruning removes 40-45% of features with zero accuracy loss",
    "Catch22 fails when class differences are local timing events
(ItalyPowerDemand)",
    "Self-supervised embeddings don't help for regular, periodic
signals (ETTh1)",
]
for i, t in enumerate(takeaways, 1):
    print(f"  {i}. {t}")
```

```
=================================================================
  CSCE 5222 – Group 18 – Complete Results Summary
=================================================================

TASK 1: Classification (UCR Datasets)
  Dataset                   Raw Acc   C22 Acc   SHAP Acc  Winner
  ———————————————————————————————————————————————————————————————
  GunPoint                   0.8267    0.8267     0.8267  Raw
  ECG200                     0.7400    0.7900     0.7800  Catch22
  ItalyPowerDemand           0.9650    0.7862     0.7804  Raw
  SyntheticControl           0.8567    0.9400     0.9500  SHAP
  TwoLeadECG                 0.7937    0.8815     0.8534  Catch22

TASK 2: Forecasting (ETTh1 – Oil Temperature)
```

```
  Method                                MSE       MAE
  ──────────────────────────────────────────────────
  Raw (lag-1 baseline)            1.2856    0.7939
  Lag + Rolling                   0.4836    0.5118
  Lag + Rolling + Embed           0.4807    0.5243
  Lag + Rolling + SHAP            0.4734    0.5090

TASK 3: Anomaly Detection (WESAD-like PPG)
  Method                     F1-Score   ROC-AUC
  ──────────────────────────────────────────────
  Raw Windows                  0.8710    0.9959
  Catch22                      0.6129    0.9363
  Catch22 + SHAP               0.6452    0.9439

====================================================================
  Key Takeaways
====================================================================
  1. Catch22 beats raw features on 3/5 UCR datasets (SyntheticControl
+6.3%)
  2. Lag+Rolling reduces forecasting MSE by ~82% over a lag-1 baseline
  3. SHAP pruning removes 40-45% of features with zero accuracy loss
  4. Catch22 fails when class differences are local timing events
(ItalyPowerDemand)
  5. Self-supervised embeddings don't help for regular, periodic
signals (ETTh1)
```

CSCE 5222 Feature Engineering | Group 18 | University of North Texas
Ashish Rathnakar Shetty (ID: 11808466) · Kushal Sai Venigalla (ID: 11852559)